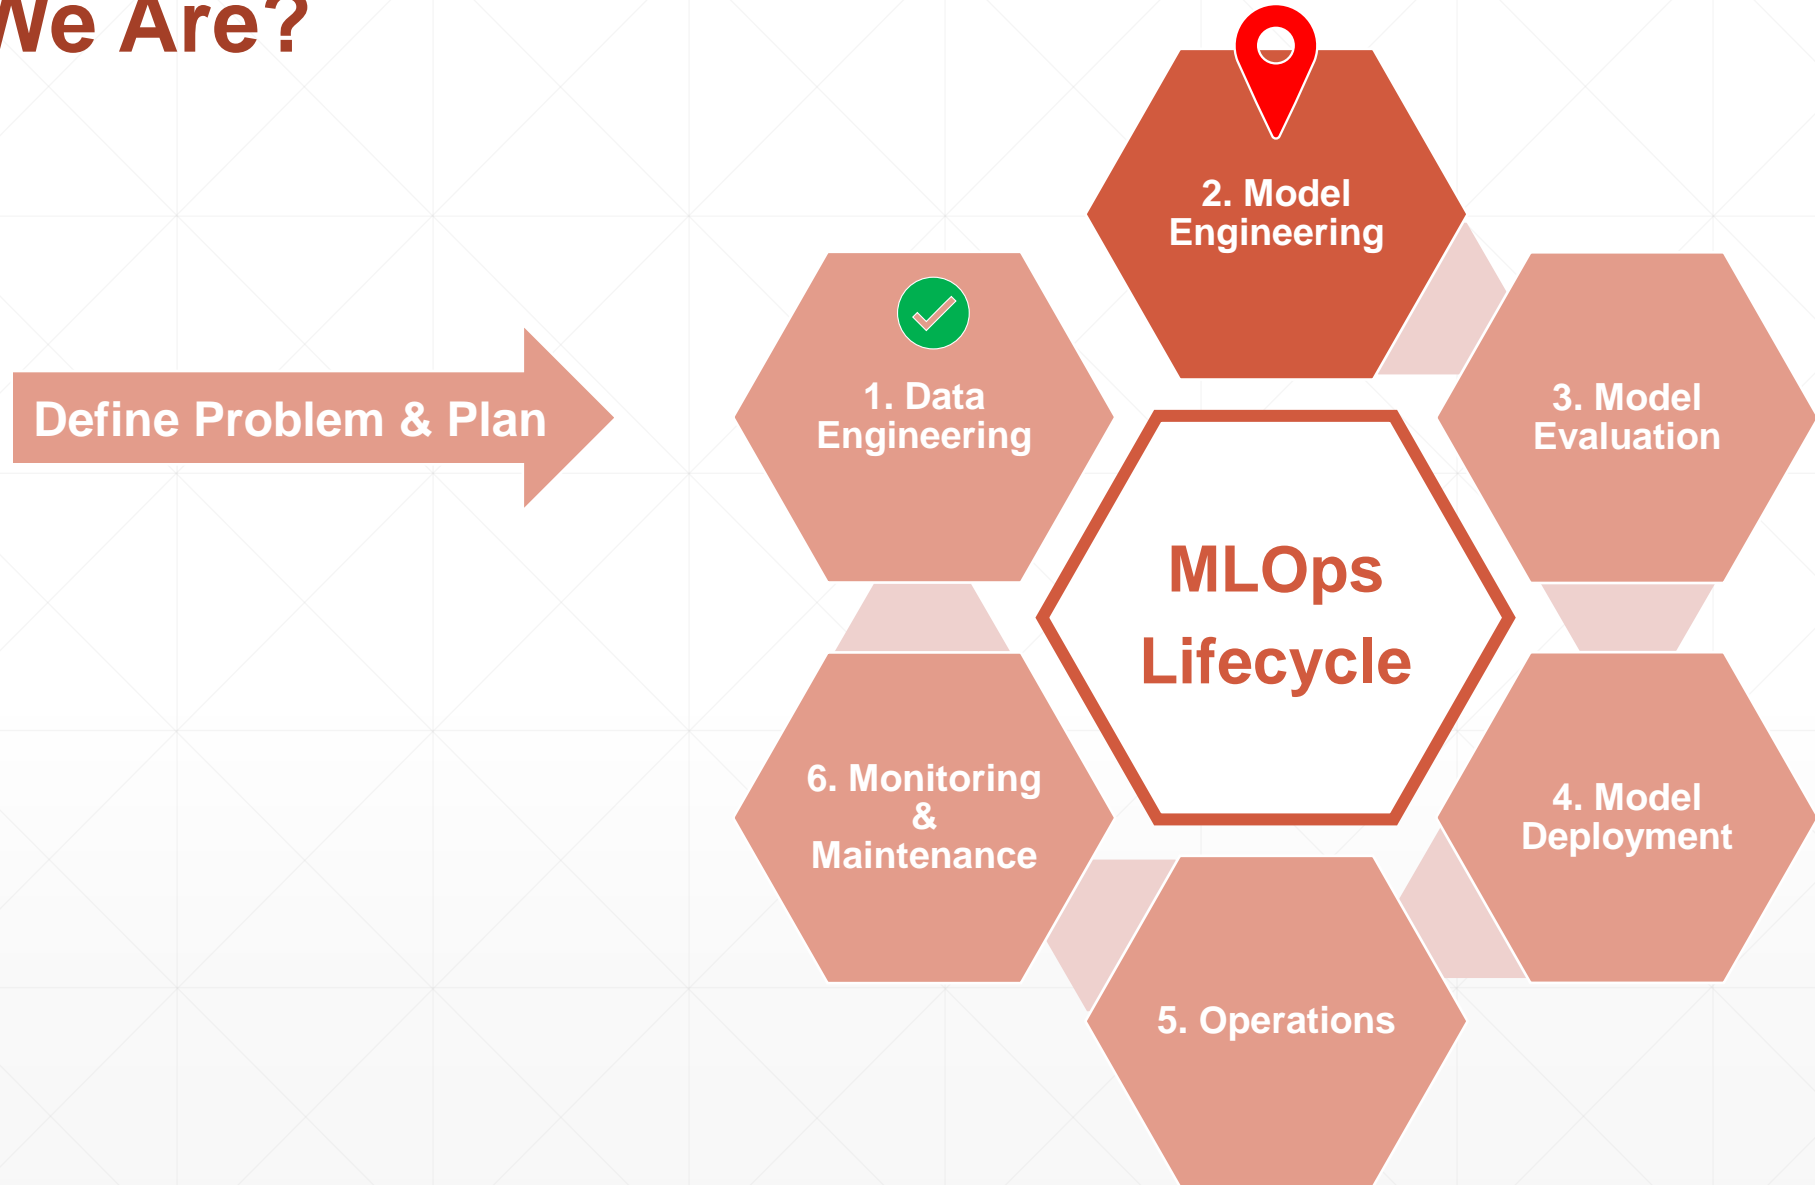


# Model Development in MLOps

---

Understand and Implement Production-Grade Machine Learning Operations

# Where We Are?



# Recap 1/

## **Stage 1: Business Understanding & Problem Definition**

- Identify the business problem.
  - Define success metrics (e.g., accuracy, precision, recall).
  - Understand constraints (data availability, computational resources).
-

## Recap 2/

### Stage 2: Data Engineering & Feature Engineering

- **Data Collection:** Gather relevant data from various sources (databases, APIs, logs).
  - **Data Preprocessing:** Handle missing values, outliers, and inconsistencies.
  - **Feature Engineering:** Transform raw data into meaningful features.
  - **Data Versioning:** Use tools like **DVC**, **LakeFS**, or **MLflow** for data versioning.
-

## Recap 3/

### Stage 3: Model Development & Experimentation

- **Select Model Architecture:** Choose between traditional ML models (e.g., Random Forest, XGBoost) or deep learning models (CNNs, RNNs).
  - **Hyperparameter Tuning:** Optimize parameters using techniques like grid search, random search, or Bayesian optimization.
  - **Logging Experiments:** Use **MLflow**, **Weights & Biases**, or **TensorBoard** to track experiments.
  - **Code Versioning:** Store code in **GitHub**, **GitLab**, or **Bitbucket**.
-

## Recap 4/

### Stage 4: Model Training & Evaluation

- **Train Model:** Use GPUs or TPUs for faster training.
  - **Evaluate Performance:** Compute metrics like accuracy, RMSE, F1-score.
  - **A/B Testing:** Compare different models on a validation set.
  - **Bias & Fairness Testing:** Ensure the model does not introduce bias.
-

## Recap 5/

### Stage 5: Model Packaging & Versioning

- **Convert Model:** Save models in formats like **ONNX**, **TF SavedModel**, or **MLflow Model Format**.
  - **Version Control:** Use **MLflow Model Registry** or **DVC** for model versioning.
  - **Containerization:** Package models using **Docker** for portability.
-

## Recap 6/

### Stage 6: Model Deployment

- **Deploy as REST API:** Use **FastAPI**, **Flask**, or **TensorFlow Serving**.
  - **Deploy to Cloud:** Use **AWS SageMaker**, **Azure ML**, or **Google Vertex AI**.
  - **Deploy with Kubernetes:** Use **Kubeflow Serving** for scalable deployment.
-



## Recap 7/

### Stage 7: Model Monitoring & Retraining

- **Monitor Performance:** Use **Prometheus**, **Grafana**, or **EvidentlyAI**.
  - **Detect Data Drift:** Identify changes in data distribution.
  - **Automate Retraining:** Set up pipelines for periodic model retraining.
-

# Learning Objectives

- MLOps Model Development
  - Reproducibility of Experiments
  - Hands-On Practice with MLflow for Experiment Tracking
  - Use Git for Code Versioning
  - Ensure Data Versioning Using Mlflow / DVC
  - Best Practices
-

# Difficulties of machine learning – The “Why”

## **Complexity of the ML Lifecycle**

The machine learning process involves multiple stages, from data preprocessing to model deployment and monitoring. Each stage requires careful management to ensure efficiency and reproducibility

## **Experiment Management**

Data scientists often conduct numerous experiments with different parameters, datasets, and algorithms. Without a centralized system, tracking these experiments and their outcomes becomes cumbersome.

## **Reproducibility**

Ensuring that experiments yield consistent results across different environments and runs is crucial. This involves managing code versions, parameters, and library dependencies.

---

# Difficulties of machine learning – The “Why”

## Deployment Consistency

With many ML libraries available, deploying models consistently can be challenging. MLflow standardizes this process, ensuring models are packaged and deployed reliably.

## Model Management

As teams produce multiple models, managing their lifecycle—versioning, testing, and deployment—becomes complex without a centralized platform

---

# Why Manage Experiments?

Managing experiments is crucial for several reasons:

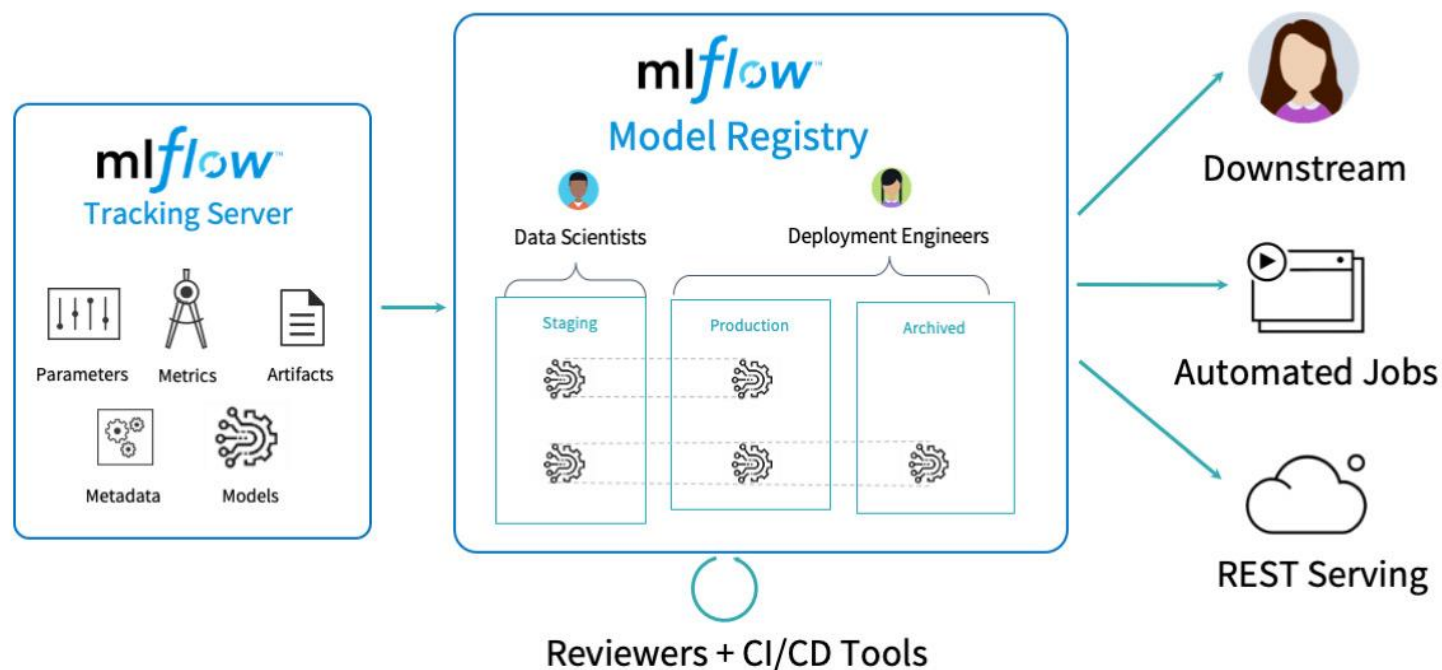
1. **Transparency and Comparison:** By logging experiments, data scientists can compare different models' performance and identify the best approach.
  2. **Reproducibility:** Experiment management ensures that results are consistent across runs, which is vital for validating findings and scaling models.
  3. **Collaboration:** Centralized experiment tracking facilitates collaboration among team members by providing a clear overview of all experiments and their outcomes.
  4. **Efficiency:** Managing experiments reduces the time spent on trial-and-error approaches and helps in optimizing resources by focusing on promising models.
-

# Problems That Cause Data Scientists to Need MLflow

- **Lack of Standardization:** Without a standardized tool, teams often work with different model versions stored in various locations, making it difficult to compare performance and manage models effectively.
  - **Inefficient Collaboration:** Collaboration is hindered when team members cannot easily access or reproduce each other's work due to inconsistent environments or missing documentation.
  - **Difficulty in Scaling:** As projects grow, managing and deploying models becomes increasingly complex without a scalable platform like MLflow.
  - **Risk of Errors:** Manual tracking and deployment processes are prone to errors, which can lead to incorrect conclusions or model failures in production.
-

# Introduction to MLflow

MLflow is an open-source platform created by Databricks to streamline the machine learning lifecycle. It provides tools for managing experiments, packaging code into reproducible runs, and deploying models in a consistent, scalable, and easy-to-monitor manner.



# Core Components of MLflow

MLflow consists of several core components that work together to manage the ML workflow:

## mlflow™ Components

mlflow™

Tracking

mlflow™

Projects

mlflow™

Models

mlflow™

Model Registry



# Core Components of Mlflow - Tracking

## 1. MLflow Tracking

- **Purposes:**
    - Logs and tracks machine learning experiments, including parameters, metrics, and artifacts.
    - Query data from experiment.
    - Store models, artifacts and code.
  - **Features:** Allows data scientists to visualize and compare different runs to identify the best-performing model. It can be used locally or remotely via the MLflow server.
  - **APIs:** Supports Java, Python, R, and REST APIs for recording and querying runs.
-

# Core Components of Mlflow - Projects

## 2. MLflow Projects

- **Purpose:**
    - Packages machine learning code into a reusable form that can be easily reproduced and shared.
  - **Features:** Each project is a directory with code or a Git repository, using a descriptor file to specify dependencies and execution methods. Projects can be chained into multi-step workflows.
  - **Example:** A project might include a conda.yaml file for specifying a Python environment.
-

# Core Components of Mlflow - Models

## 3. MLflow Models

- **Purpose:**
    - Standardize models for deployment.
    - Build customized models.
  - **Features:** Models are saved as directories with a descriptor file listing supported flavors (e.g., TensorFlow, Python function). Supports deployment to platforms like AWS SageMaker and Azure ML.
-

# Core Components of Mlflow - Model Registry

## 4. MLflow Model Registry

- **Purpose:**
    - Provides a centralized model store for managing the lifecycle of MLflow models (Store and version ML models)
    - Load and deploy ML models.
  - **Features:** Offers model lineage, versioning, stage transitions (e.g., staging to production), and annotations. Facilitates collaboration and governance.
-

# MLflow - Additional Components & Features

## 5. MLflow Deployments for LLMs

- **Purpose:** Streamlines access to SaaS and OSS Large Language Models (LLMs) through standardized APIs.
- **Features:** Enhances security with authenticated access and provides a common API interface for prominent LLMs.

## 6. Evaluate

- **Purpose:** Facilitates in-depth model analysis and comparison.
- **Features:** Supports objective comparison of traditional ML algorithms and cutting-edge LLMs.

## 7. Prompt Engineering UI

- **Purpose:** Offers a dedicated environment for prompt experimentation, refinement, evaluation, testing, and deployment.
  - **Features:** Ideal for working with LLMs and other models requiring precise input crafting.
-

# MLflow - Additional Components & Features

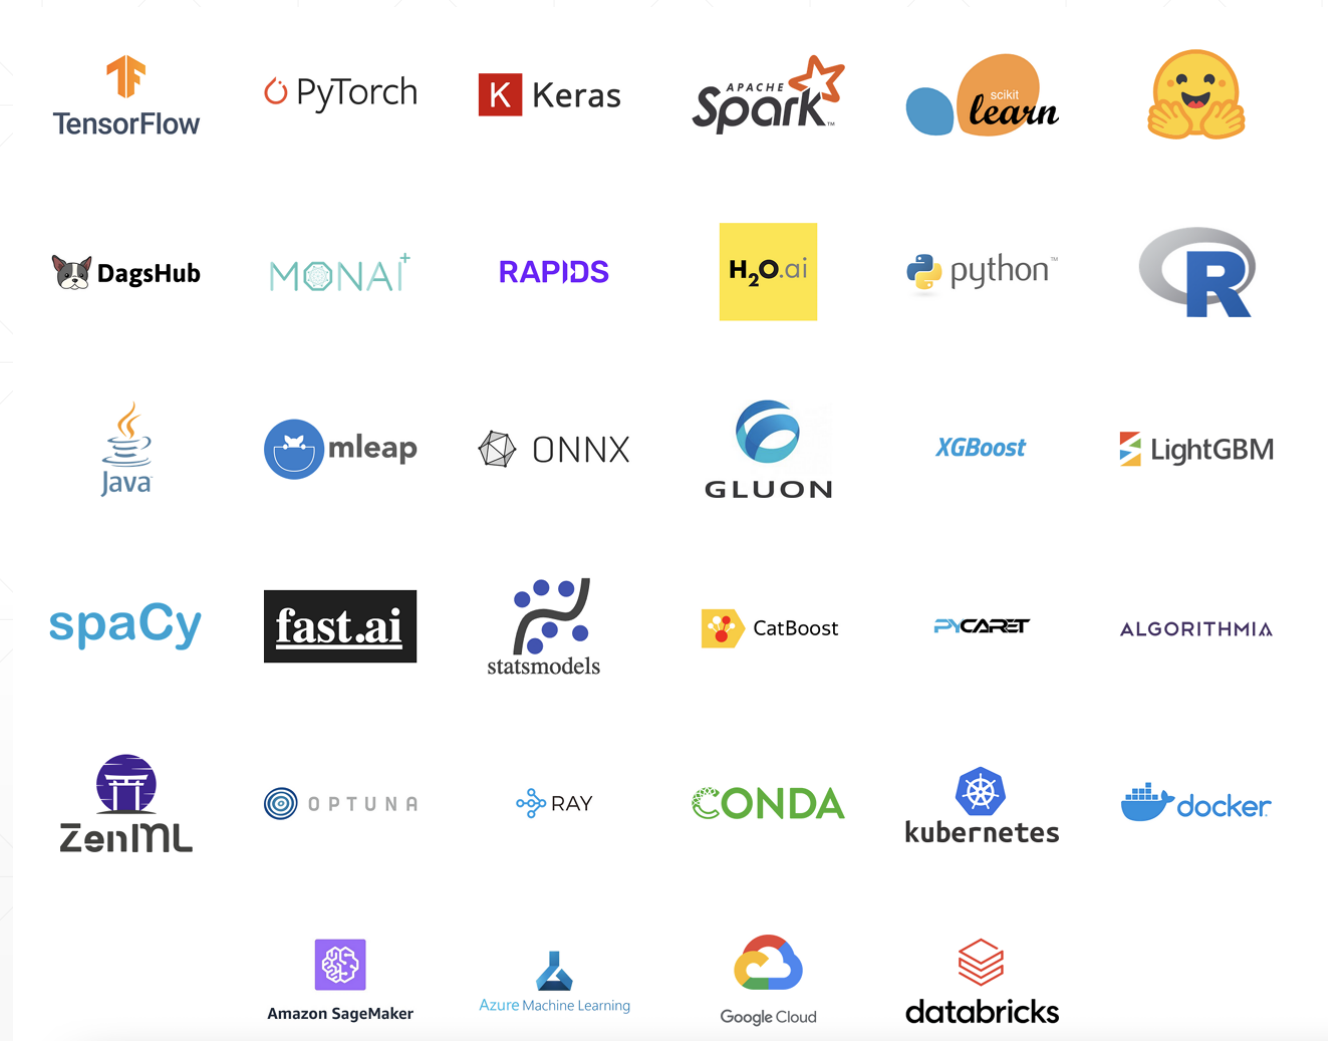
## 8. Recipes

- **Purpose:** Guides the structuring of ML projects for real-world deployment scenarios.
- **Features:** Provides recommendations to ensure functional end results.

## Advanced Features and Use Cases

- **Nested Runs:** Organize experiments in a hierarchical structure for more structured experimentation.
  - **Rich Metric Visualization:** Utilize advanced tools to analyze metrics over time and across different runs.
  - **Automated ML Pipelines:** Use predefined templates for common tasks and dynamic step execution to adapt workflows based on data or previous results.
  - **Multi-Model Endpoints:** Serve multiple models or versions from a single endpoint to optimize resource usage.
  - **Plugin Ecosystem:** Extend MLflow's functionality with custom plugins for storage, authentication, etc..
-

# Mlflow - Integrations



# Practice: Setup Environment

Execute following commands (Linux required)

1. `cd ~/dev`
  2. `mkdir mlflow-practice`
  3. `python -m venv venv`
  4. `echo mlflow > requirements.txt`
  5. `source venv/bin/activate`
  6. `pip install -r requirements.txt`
  7. `mkdir logs`
  8. `mlflow ui >> logs/log_file.txt 2>&1 &`
-



# Mlflow Experiments

The screenshot displays the Mlflow Experiments interface. The top navigation bar includes the Mlflow logo (2.10.2), tabs for 'Experiments' and 'Models', and links for 'GitHub' and 'Docs'. The 'Experiments' tab is active, showing a search bar and a list of experiments. The 'search-run-guide' experiment is selected and highlighted with a red circle. A red arrow points from this circle to the 'search-run-guide' experiment name in the main view. The main view shows the experiment details, including the Experiment ID (928782070984408674) and the Artifact Location. Below this, there is a search bar with the query 'metrics.rmse < 1 and params.model = "tree"', filters for 'Time created' and 'State: Active', and a '+ New run' button. The 'Table' view is selected, showing a list of runs. A red circle highlights the runs table, and a red arrow points from the text 'MLflow Runs' to it. The runs table has columns for Run Name, Created, Dataset, Duration, Source, and Models. The runs are listed in descending order of creation time.

**Experiments** (+) (-)

Search Experiments

☐ Default ☒ search-run-guide

**search-run-guide** Provide Feedback

Experiment ID: 928782070984408674 Artifact Location: file:///Users/michael.berk/dev/mlflow/search-query-syntax/test/mlruns/928782070984408674

> Description Edit

Q metrics.rmse < 1 and params.model = "tree" Time created State: Active

Datasets Sort: Created Columns

Table Chart Evaluation Experimental

	Run Name	Created	Dataset	Duration	Source	Models
<input type="checkbox"/>	mercurial-shark-477	39 seconds ago	also custom (4965bbbd) Test	26ms	delete.py	-
<input type="checkbox"/>	gaudy-toad-977	39 seconds ago	also custom (6efd27a7) Test	24ms	delete.py	-
<input type="checkbox"/>	puzzled-gnu-272	39 seconds ago	also custom (80b3ed44) Test	25ms	delete.py	-
<input type="checkbox"/>	enthused-ram-468	39 seconds ago	also custom (4b218f0d) Test	25ms	delete.py	-
<input type="checkbox"/>	vaunted-stork-741	40 seconds ago	also custom (cca8b58e) Test	26ms	delete.py	-
<input type="checkbox"/>	nimble-gnat-675	40 seconds ago	custom (65e99d28) Train	26ms	delete.py	-
<input type="checkbox"/>	orderly-quail-568	40 seconds ago	custom (5c432a8b) Train	26ms	delete.py	-
<input type="checkbox"/>	melodic-lynx-301	40 seconds ago	custom (bfc733bd) Train	26ms	delete.py	-

10 matching runs

localhost:5000/#/

# Working with Experiments

## MLflow Client

- Create Experiments

```
client.create_experiment("Name")
```

- Tag Experiments

```
client.set_experiment_tag("Name",  
k, v)
```

- Delete Experiments

```
client.delete_experiment("Name")
```

## MLflow module

- Create Experiments

```
mlflow.create_experiment("Name")
```

- Tag Experiments

```
mlflow.set_experiment_tag(k, v)
```

- Delete Experiments

```
mlflow.delete_experiment("Name")
```

- Set Experiment

```
mlflow.set_experiment("Name")
```

---


# Starting a New Experiment

```
import mlflow
# Create new Experiment
mlflow.create_experiment("My Experiment")
# Tag new experiment
mlflow.set_experiment_tag("scikit-learn", "lr")
# Set the experiment
mlflow.set_experiment("My Experiment")
```

---

## Let's Practice – MLflow experiments

PRACTICE



# Training runs

- How MLflow is organized
  - New run equals new model training
  - A run is placed within an experiment
  - Invoked via `mlflow.start_run()`
-

# Training runs



# Start a training run

```
import mlflow
```

```
# Start a run
```

```
mlflow.start_run()
```

```
<ActiveRun: >
```

```
# End a run
```

```
mlflow.end_run()
```

---

# Setting a training run variable

```
import mlflow
# Set experiment
mlflow.set_experiment("My Experiment")
# Start a run
run = mlflow.start_run()
# Print run info
run.info
```

```
<RunInfo: artifact_uri='./mlruns/0/9de5df4d19994546b03dce09aefb58af/artifacts',
  end_time=None, experiment_id='31', lifecycle_stage='active',
  run_id='9de5df4d19994546b03dce09aefb58af', run_name='big-owl-145',
  run_uuid='9de5df4d19994546b03dce09aefb58af', start_time=1676838126924,
  status='RUNNING', user_id='user'>
```



# Logging to MLflow Tracking

- **Metrics**

- `log_metric("accuracy", 0.90)`
- `log_metrics({"accuracy": 0.90, "loss": 0.50})`

- **Parameters**

- `log_param("n_jobs", 1)`
- `log_params({"n_jobs": 1, "fit_intercept": False})`

- **Artifacts**

- `log_artifact("file.py")`
  - `log_artifacts("./directory/")`
-

# Logging a run

```
import mlflow
# Set Experiment
mlflow.set_experiment("LR Experiment")

# Start a run
mlflow.start_run()

# Model Training Code here
lr = LogisticRegression(n_jobs=1)

# Model evaluation Code here
lr.fit(X, y)
score = lr.score(X, y)
```

```
# Log a metric
mlflow.log_metric("score", score)

# Log a parameter
mlflow.log_param("n_jobs", 1)

# Log an artifact
mlflow.log_artifact("train_code.py")
```

# Searching runs

```
mlflow.search_runs()
```



# Searching runs – Filter searches

- `max_results` - maximum number of results to return.
  - `order_by` - column(s) to sort in `ASC` ending or `DESC` ending order.
  - `filter_string` - string based query.
  - `experiment_names` - name(s) of experiments to query.
-

# Searching runs – examples

```
import mlflow
# Filter string
f1_score_filter = "metrics.f1_score > 0.60"
# Search runs
mlflow.search_runs(experiment_names=["Insurance Experiment"],
                  filter_string=f1_score_filter,
                  order_by=["metrics.precision_score DESC"])
```

---

## Let's Practice – MLflow tracking runs

# PRACTICE



[thehn/getting\\_started\\_mlflow](#)